

Script Assembler – Generation of Test Scripts for KeTop Devices

DIPLOMA THESIS

submitted for the

Reife- und Diplomprüfung

at the

Department of Informatik

Submitted by:

Filip Schauer
Axel Csomány
Adam Höllerl

Supervisor:

Adrian Kern

Project Partner:

Martin Wieser, KEBA Group AG

Leonding, April 7, 2026

I hereby declare that I have composed the presented paper independently and on my own, without any other resources than the ones indicated. All thoughts taken directly or indirectly from external sources are properly denoted as such.

This paper has neither been previously submitted to another authority nor has it been published yet.

The presented paper is identical to the electronically transmitted document.

Leonding, April 7, 2026

Filip Schauer , Axel Csomány & Adam Höllerl

Abstract

Brief summary of our amazing work. In English. This is the only time we have to include a picture within the text. The picture should somehow represent your thesis. This is untypical for scientific work but required by the powers that are. Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.



Zusammenfassung

Zusammenfassung unserer genialen Arbeit. Auf Deutsch. Das ist das einzige Mal, dass eine Grafik in den Textfluss eingebunden wird. Die gewählte Grafik soll irgendwie eure Arbeit repräsentieren. Das ist ungewöhnlich für eine wissenschaftliche Arbeit aber eine Anforderung der Obrigkeit. *Bitte auf keinen Fall mit der Zusammenfassung verwechseln, die den Abschluss der Arbeit bildet!* Suspendisse vel felis.

Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.



Contents

I. Introduction [X, F]	1
1. Background [X]	2
1.1. Testing in Industrial Systems	2
1.2. Company Overview: KEBA Group AG	2
2. Initial Problem [X]	4
2.1. Current Testing Strategy and its Flaws	4
2.2. Goals and Objectives	4
3. Used Technologies [F]	6
3.1. Python	6
3.2. Robot Framework	7
3.3. Monaco Editor	7
3.4. PostgreSQL	7
II. Keywords [F, X, A]	9
4. Introduction to Keywords [F]	10
4.1. Robot files	10
4.2. What Keywords Are	11
4.3. Examples	12
4.4. Recursion	13
4.5. Concatenation and Output	14
5. The Different Keywords [X]	15
5.1. Options	15
5.2. Modifiers	16
5.3. Visual Example	18

6. Why This Syntax? [A]	20
 III. EBNF [A, F]	 21
7. Overview of EBNF [A]	22
7.1. Motivation for Formal Syntax Descriptions	22
7.2. The Backus-Naur Form (BNF)	23
7.3. The Evolution to Extended BNF (EBNF)	23
7.4. Wirth's Approach to Syntax Definition	24
7.5. Core Syntactic Extensions	26
7.6. ISO/IEC 14977 Standards for Assembler Logic	29
 8. Approaches to EBNF [A]	 33
8.1. Requirements	33
8.2. Minimal EBNF Approach	34
8.3. Detailed EBNF Approach	35
8.4. Conclusion	36
 9. A Problem with EBNF [F]	 38
 10. Our EBNF [F]	 39
 IV. Implementation [A, F, X]	 40
11. Preprocessor [A]	41
11.1. Introduction	41
11.2. Design and Transformation Rules	41
11.3. Evolution of the Algorithm	43
11.4. Implementation Details	46
11.5. Summary	50
 12. Parser [F, A]	 51
 13. Syntax Highlighting [A]	 52
13.1. Introduction	52
13.2. Monaco Editor Setup	53
13.3. Monarch Lexer for Custom Annotations	55

13.4. Troubleshooting and Error Handling	57
13.5. Limitations	59
14. Extensibility [F]	61
15. Database [X]	62
15.1. Types of Databases	62
15.2. Relational Databases	63
15.3. The TestRail Database	66
15.4. PostgreSQL	67
15.5. First Approaches	69
15.6. Database Connections in the System	69
15.7. Integration of the DBSession Class	71
15.8. The PostgreSQL Database	72
15.9. Potential Future Steps	73
Glossary	VII
References	VIII
List of Figures	XI
List of Tables	XII
List of Listings	XIII
Appendix	XIV

Part I.

Introduction [X, F]

1. Background [X]

1.1. Testing in Industrial Systems

The testing of systems is a crucial part of the development process. It not only ensures that the system functions as intended but also helps identify flaws in areas not considered during the design phase. Testing can be performed at various stages of development, from unit testing to integration testing and system testing.

In the context of industrial automation, testing is particularly important due to the complexity and critical nature of the components involved. The systems in question often control machinery and processes that can bring about significant consequences if they fail. Therefore, a robust testing strategy is essential in order to ensure the reliability of these systems.

Companies operating in this field must develop and maintain thorough testing strategies to ensure system reliability, safety and with that the quality of their products and the satisfaction of their customers. This project focuses on improving the existing testing systems of the KEBA Group AG.

1.2. Company Overview: KEBA Group AG

KEBA Group AG is an austrian technology company specializing in automation systems for industry, banking and energy applications. It designs and produces hardware and software solutions such as machine and robot controls, automated handover terminals, and electric vehicle charging stations.

For more than 50 years, they have been providing various customers with solutions to their unique problems. Among its three main branches, the Industrial Automation branch focuses primarily on the needs of the individual situation as the hardware and its software get built.

The company was founded in 1968. Today with 28 establishments in 16 different countries and more than 2000 employees they are internationally active, serving global markets from its headquarters in Linz, Austria.

Within the industrial context, this project was conducted in collaboration with KEBA Group AG.

2. Initial Problem [X]

2.1. Current Testing Strategy and its Flaws

So called "KeTops" are used as operating devices to control machines of varying structure and complexity. These machines and their corresponding KeTops differ in their configuration from each customer which is why fitting test scripts are required. These test scripts examine all aspects of the device from the installed operating system to software specific behaviors.

The "Script Assembler" which is used in current production is able to merge a number of test scripts together for this use. The hard-coded nature and resulting limited maintainability of the Script Assembler is its primary flaw. Furthermore all data concerning tests are saved and maintained in a TestRail database which although functional, heavily limits the expandability and independence of the testing environment.

2.2. Goals and Objectives

This project aims to eliminate all of the listed problems and concerns while building upon and improving the existing capabilities of the Script Assembler.

- **EBNF parsing:**

The existing structure of the test scripts should be formed into its own EBNF language which is then integrated into the Script Assembler. This eliminates a majority of the hard-coded parts of the software. Furthermore the editor should provide appropriate syntax highlighting when writing or editing a test script.

- **Keyword support:**

New functionalities should be implemented which allows the editor of test scripts to specify aspects like the version and the operating system the test has to be run on. The structure of keywords need to be precisely set in the EBNF language.

- **Database mock:**

The structure of the currently used TestRail database should be studied and an accurate mockup should be created. A seamless switch between the mock and the real TestRail instance has to be implemented as well.

3. Used Technologies [F]

Just like the TRP project for which this project aims to provide an extension for, it was mainly written in the Python programming language. Aside from that, other tools have also been used or are otherwise important to be mentioned. The goal of this section is to give a broad overview of some of the most important technologies that each make up an integral part of our final project.

3.1. Python

Python is a high-level, interpreted, dynamically typed object oriented programming language with a focus on allowing programmers to create applications, scripts, or connections between existing programs, quickly.[1]

3.1.1. Modules

Importantly, Python allows for the import of “modules”, which are essentially files containing Python code which allow for importing functions, classes, or similar, into other Python files.[2]

While it is possible for modules to be local files, they are also commonly downloaded from the internet using “package managers”. An commonly used example of a package manager for Python, which is also the only Python package manager used for this project, is known as `pip`. [3]

The external modules used in this project include, but are not limited to:

- `lark`, which allows for parsing of a given piece of text using a custom grammar language which is based on EBNF. The core of our parser is built around this library.[4]

- **semantic-version**, which is a small library that provides tools to handle “SemVer”.^[5] “SemVer” stands for “Semantic Versioning”, and it refers to a scheme that intends to dictate how a software’s version number is to be assigned and incremented.^[6] This library is used to for parsing version strings for our “@VERSION Keyword”.
- **psycopg**, which allows for communication with a “PostgreSQL” database.^[7] This library is used for our “Mockup Database” system.

3.2. Robot Framework

Robot Framework is a framework for test automation. It allows for defining so-called “Test Suites”, which are files ending in `.robot` and may contain any amount of test cases.^[8]

In the context of KEBA Group AG, Robot Framework is used to execute various test cases on KeTop devices. However, since our parser only acts as a preprocessor to robot script files, treating them as arbitrary text to be either included, excluded, and/or moved to different positions depending on the defined keywords and context, Robot Framework’s actual syntax is, for the purposes of our project, irrelevant. Exceptions to this include the question of which characters of a standard “ASCII” format are syntactically relevant, and its use of “significant indentation”.

3.3. Monaco Editor

The Monaco Editor is a code editor with support for features such as “IntelliSense” and “Syntax Colorization”, which also makes up the core of the commonly used “VS Code” application.^[9]

Our project makes use of its “Syntax Colorization” feature, which we instead refer to as “Syntax Highlighting” in this thesis.

3.4. PostgreSQL

PostgreSQL is an “ACID”-compliant object relational database management system which uses an extended version of the “SQL” language.^[10]

PostgreSQL was our choice for our “Mockup Database” system, which is needed to allow for running our project independently from the proprietary “TestRail” test management software.

Part II.

Keywords [F, X, A]

4. Introduction to Keywords [F]

In programming, a “keyword” is a predefined word with a reserved meaning. Keywords can generally not be used for any other purpose, and they make up a programming language’s syntax.[11]

In our Script Assembler, we use custom keywords in a similar way, although it may be difficult to define whether or not it is useful to refer to our parser’s syntax as a “programming language”.

In this chapter, the term “whitespace” is defined to refer to any amount of space or tab characters. Here, for convenience, line-breaks are chosen to specifically not fall under said definition.

4.1. Robot files

In the TRP project, the Robot Framework provides the primary scripting language that makes up the tests to be run on KeTop devices.

A Robot file contains so-called “test cases”, which are blocks commonly made up of a few lines of code.[8] Despite most parts of the syntax being irrelevant to the Script Assembler’s workings, in order to get a feel for the nature of this project it may still be useful to have a look at an example of a Robot script file:

```
1  *** Settings ***
2  Documentation      A test suite for valid login.
3  ...
4  ...               Keywords are imported from the resource file
5  Resource           keywords.resource
6  Default Tags       positive
7
8  *** Test Cases ***
9  Login User with Password
10     Connect to Server
11     Login User      ironman      1234567890
12     Verify Valid Login    Tony Stark
13     [Teardown]    Close Server Connection
14
15  Denied Login with Wrong Password
16     [Tags]    negative
17     Connect to Server
18     Run Keyword And Expect Error    *Invalid Password    Login User    ironman
19     123
```

```

19     Verify Unauthorised Access
20     [Teardown]      Close Server Connection

```

[8]

As is made clear from the example above, some lines may be preceded by whitespace. Since the Robot Framework achieves nesting by having lines be preceded by an integer multiple of exactly four spaces certain lines of code may be syntactically erroneous or interpreted differently when that amount of whitespace is changed.[12] This is an example of a programming convention that in this thesis will be referred to as “significant indentation”. Python, for instance, works in a similar way, requiring a predefined use of leading whitespace to achieve grouping of statements.[13]

Despite Robot Framework’s general ease of use, proficient support for extensibility, and versatile testing capabilities, the control over the specific execution flow is limited.[14] In the context of KEBA Group AG, a single KeTop may need to run multiple Robot script files, and a Robot script file may need to be reused for multiple KeTops. However, since it may be required for a test to be made of a single script file, it might be necessary to “stitch together”, or in other words to concatenate, the contents of multiple Robot script files into one.

Without Script Assembler, a user may run into issues where they, for instance, would like a certain piece of code to only be included on KeTops satisfying a certain condition like being of a predefined operating system, or for certain pieces of code to be executed at a time that is either before or after the remaining parts, as the Robot Framework may not provide a useful way to achieve the required flow.

4.2. What Keywords Are

Our project allows for custom keywords, all of which being preceded by the `@` character, to be defined to act as either “options” or “modifiers”. In a provided file, these keywords must be defined at the start of the line, optionally preceded by indentation for nesting purposes. Lines containing such valid keywords will not be included in the concatenated output file.

Each of the keywords, no matter if option or modifier, may be defined to either have no arguments, or have optional or required arguments. If set, these arguments may

then consist of any sequence of characters without line-breaks that is prefixed by the keyword with at least one whitespace inbetween.

The difference between options and modifiers lies in that options do not include “body text”, while modifiers do. “Body text” is defined to be made up of all lines of text that, when read from top to bottom, succeed the line of the keyword, up until and not including the first line that is found that does not include at least one additional level of indentation than that of the line of the keyword, or up until the end of the file, ignoring all lines that are either empty or are solely made of whitespace. A level of indentation is defined to be made up of either four spaces or one tab character.

To make up for the extra level of indentation used to define the span of the body text, the parsed body text will be stripped of that one line of indentation. Non-keyword lines, and lines that do not belong to the block body of a modifier will simply be ignored by the Script Assembler and will thus be appended into the output Robot file as-is. This means that, if the provided file does not contain anything that may be matched as a keyword, which, in the case of a valid Robot file, may only be the case in rare exceptions, as the @ character is generally not utilized in Robot files that are to be executed on a KeTop.

4.3. Examples

4.3.1. Options

For instance, take the following code snippet, which includes the option @NO_KETOP, which takes no arguments:

```
1  @NO_KETOP
2  Install MS
```

When the line containing @NO_KETOP is encountered, the `no_ketop` flag, for which the concrete behaviour is explained in the following chapter, is enabled for the output script. Since that specific line will not be included in the output script, assuming the snippet above to be the only input to the parser, the output file may look like the following:

```
1  Install MS
```

4.3.2. Modifiers

Another example is using the modifier `@VERSION`:

```

1  Execute on Ketop      echo A
2  @VERSION BIOS > 1.0.1
3      IF      ${SINGLE_MODE} is $False
4          Execute on Ketop      echo B
5      END
6  Execute on Ketop      echo C

```

Here, the line two, which is made up of `@VERSION BIOS > 1.0.1`, represents a boolean condition relating to the version of the KeTop device. The block body of this condition is made up by the lines three to five. According to the specification of the `@VERSION` keyword, that block body will be included only if the condition stated in its argument is satisfied.

This means that, again assuming the snippet above to be the only input to the parser, if the condition `BIOS > 1.0.1` is satisfied, the output may look like the following:

```

1  Execute on Ketop      echo A
2  IF      ${SINGLE_MODE} is $False
3      Execute on Ketop      echo B
4  END
5  Execute on Ketop      echo C

```

Note the removal of one level of indentation of what used to be the block body. If, instead, the case of the condition is not satisfied, the output may instead be:

```

1  Execute on Ketop      echo A
2  Execute on Ketop      echo C

```

4.4. Recursion

Keywords may also be defined recursively. Recursion is generally defined as being “the process in which a function calls itself directly or indirectly”[15]. In the case of Script Assembler, it refers to the fact that block bodies of modifiers may also include modifiers, of which the block bodies may then also include modifiers, and so forth. To show this possibility, consider the following example:

```

1  @BRACKETSTART
2      @OS WIN10
3          @VERSION BIOS >= 1.7.0
4              Upload Script To KeTop      E2Speed.ps1

```

Here, the block body of `@BRACKETSTART` is made up of lines two to four, of which the block body of `@OS` is made up of lines three and four, of which the block body of `@VERSION` is made up of line four. Despite the order of “which blocks are nested inside which” not making a difference with how all existing modifiers are specified, in general,

the parser will interpret this from the outside to the innermost level of nesting, similar to how nesting is handled by most ordinary programming languages.

This means that, if the above snippet is the only input file that contains the `@bRACKETSTART` keyword, given how that keyword is specified, the part of the output script that is the so-called “bracket start” will be:

```
1 Upload Script To KeTop      E2Speed.ps1
```

in the case of both of the conditions from the respective specifications of the `@OS` and `@VERSION` keywords and with each of their respective arguments being met. Otherwise, the “bracket start” of the output file will be empty.

4.5. Concatenation and Output

In general, the files get concatenated in the same order that they are given, unless keywords that change that predefined flow are used. Then, to assure the final output Robot file is valid and executable on a KeTop, the following template code will be appended before the output from the entire parsing and concatenation operations:

```
1 # Template file for HMI Automation
2 *** Settings ***
3 Resource      ../../../../Robot/Resources/keywords.resource
4 Resource      ../../../../Robot/Resources/variables.resource
5 Suite Setup   KeTop Suite Setup
6 Suite Teardown KeTop Suite Teardown
7 Test Setup    KeTop Test Setup
8 Test Teardown KeTop Test Teardown
9
10 *** Test Cases ***
11 # Here the test cases are inserted
```

Finally, the resulting code will be written to a file which will then be further handled to be run either locally or on a KeTop.

5. The Different Keywords [X]

This section aims to provide a comprehensive list of keywords and their corresponding definitions and uses in the context of the Script Assembler.

5.1. Options

Options are one of the two types of keywords that can be used within a script. They are used to set specific parameters or configurations for the specific script execution, essentially acting as flags that modify certain behaviors when dealing with the script. It doesn't require a value to be assigned to it, so its presence alone is enough to trigger the intended effect. The position of the option keyword within the script is not crucial, so it can be placed anywhere within the script while keeping its intended effect intact.

During the making of the project, one option keyword has been implemented:

5.1.1. Option: NO_KETOP

The NO_KETOP option keyword can be used by typing the following line anywhere within the script:

```
@NO_KETOP
```

If this option keyword is present, the "no_ketop" flag will be set for the entire script. This ensures that certain components that check if a KeTop device is connected are disabled, allowing the script to run without the prerequisite of a KeTop device. This is most useful for specific test scripts that have to be run without a KeTop, like installation scripts.

5.2. Modifiers

Modifiers are the second type of keywords that can be used within a test script. They differ from options in that they require a body, specified by the presence of indentation, to function properly. The body of a modifier contains further lines of the test script, the modifier keyword itself defines under what conditions the body should be included in the final script. Generally speaking, modifiers are used to either conditionally include or exclude certain parts of the script, or set them in a specific order.

Bodies of modifiers can also contain other modifiers, allowing the creation of complex and more specific test scripts.

During the making of the project, the following modifier keywords have been implemented:

- VERSION
- OS
- BRACKETSTART and BRACKETEND
- PREAMBLE

5.2.1. Modifier: VERSION

The VERSION modifier keyword can be used following the syntax:

```
@VERSION <component> <comparison> <version>  
    <body>
```

Where:

- **<component>** is the label of the component that is being checked, for example "BIOS".
- **<comparison>** is a comparison operator, such as ">", "<", ">=", "<=", "==" or "!=".
- **<version>** is the version of the component that the test script is being run on, in the format of "X.Y.Z" where X, Y and Z are integers.

If the version modifier keyword is present, the body of the modifier will only be included in the final script if the condition concerning the component version is met. This allows

for the creation of test scripts that behave differently based on the component version of the device.

5.2.2. Modifier: OS

The OS modifier keyword can be used following the syntax:

```
@OS <os>  
    <body>
```

Where:

- **<os>** is the label of the operating system that is being checked, for example "WIN10".

If the OS modifier keyword is present, the body of the modifier will only be included in the final script if the device's operating system matches the specified system.

5.2.3. Modifier: BRACKETSTART and BRACKETEND

The BRACKETSTART and BRACKETEND modifier keywords do not require any parameters, only the keyword itself and its body:

```
@BRACKETSTART  
    <body>
```

```
@BRACKETEND  
    <body>
```

If a script containing a BRACKETSTART or BRACKETEND modifier keyword is being run together with other test scripts, the body of these modifiers will be run before and after the content of all other test scripts respectively. If a batch of test scripts contains multiple scripts containing these modifiers, the bodies of the modifiers will be run in the order they appear in the batch, meaning that the body of the first BRACKETSTART modifier will be run before the body of the second BRACKETSTART modifier. The same principle applies to the BRACKETEND modifiers.

Therefore, these modifiers can be used to set up tests that require a certain amount of time to have passed, before the result can be observed. For example, if a test script is meant to check the behavior of a device after it has been active for 30 minutes, the

BRACKETSTART modifier can be used to set up the test, and the BRACKETEND modifier can be used to check the results after the 30 minutes have passed. In the meantime, other test scripts can be run between the two "bracket tests", allowing for a more efficient use of time when running a batch of test scripts.

5.2.4. Modifier: PREAMBLE

The PREAMBLE modifier keyword can be used following the syntax:

```
@PREAMBLE <index>
    <body>
```

Where:

- **<index>** is an integer that specifies the order in which the preambles are run, with 1 being the first preamble to be run.

When one or multiple test scripts containing a PREAMBLE modifier keyword is being run together with other test scripts, the body of the PREAMBLE keyword will be run before the content of all other test scripts, including lines in the body of a BRACKETSTART modifier. In the case of multiple PREAMBLE modifiers being present in the batch of test scripts, the bodies will be run in the order of their specified index.

The PREAMBLE modifier can be used to write scripts that test certain installation processes and their successful completion before the rest of the test scripts are run.

5.3. Visual Example

Consider the following diagram:

Here, we can see how some specific test scripts are arranged in the final script.

1. The bodies of the PREAMBLE modifiers are ordered first if present. In the case of multiple PREAMBLE scripts being present, the specified index is used to order them accordingly. As seen in the diagram, leaving out certain indices doesn't lead to any issues, the ordering occurs as intended.




Figure 1.: Simple Diagram depicting the order scripts with a certain modifier keyword are run in.

2. The bodies of the BRACKETSTART modifiers are ordered next if present. In the case of multiple BRACKETSTART scripts being present, the order they appear in the batch is maintained.
3. At this point, the content of all test scripts not containing any order-modifying keywords are placed. Their order is determined by the order they appear in the batch.
4. Lastly, the bodies of the BRACKETEND modifiers are ordered if present. In the case of multiple BRACKETEND scripts being present, they follow the same principle as the BRACKETSTART modifiers, meaning that the order they appear in the batch is maintained.

6. Why This Syntax? [A]

Part III.

EBNF [A, F]

7. Overview of EBNF [A]

7.1. Motivation for Formal Syntax Descriptions

To be able to run tests automated on every kind of KeTop, there is a need to modify test scripts depending on the given environment. That is why an extended Robot Framework language was needed. The Extended Backus-Naur Form (EBNF) was chosen for this, to make this new language easily extendable and deterministic to implement if adopted by other developers.

In practice, in the field of computer science, it is clear that a precise description of a language is the most important foundation. It is therefore essential that applications, communication protocols or programs function according to clear syntactic rules so that they are always interpreted in the same way by humans and machines. Natural language provides many different ways of describing something. There are ambiguities in language, and these can lead to errors in interpretation. These findings are not new. They have been known for a long time. Nevertheless, there is still no common standard that everyone adheres to, especially when languages have been implemented independently by different developers (cf. [16, p. V]).

Syntactic metalanguages exist to solve this problem. The goal of a metalanguage is to describe another language. This means there is a difference between the language that is described and the language used to describe it.

In programming, a metalanguage is a language used to describe the syntax and structure of another language, called the object language. For example, BNF and EBNF are metalanguages used to define programming languages such as ALGOL. The structure of sentences and expressions in a metalanguage can be defined using a so-called metasyntax. For example, one could indicate that the word “noun” can itself function as a noun in a sentence by writing: “noun” is a `<noun>`. (cf. [17])

7.2. The Backus-Naur Form (BNF)

One of the most important metalanguages for programming has been BNF (the so-called Backus-Naur Form). BNF started with ALGOL 60 in 1960. ALGOL 60 needed a clear, formal way to describe its syntax. John Backus and Peter Naur therefore created BNF. It made it possible to define language rules precisely. This was the first time syntax was described formally. The history of ALGOL 60 is crucial because it gave BNF its first real use. BNF became the foundation for defining many programming languages. It is still essential in compilers and language design today. (cf. [18, p. 3])

As BNF was first used for ALGOL 60 and became the model for many later languages, over time, people added changes or extensions, which created some confusion. Despite this, BNF showed the importance of clear, formal syntax.

Therefore, Extended BNF was later developed from BNF and included the most common extensions. The large amount of variants motivated the standardization effort that resulted in ISO/IEC 14977. (cf. [16, p. V])

7.3. The Evolution to Extended BNF (EBNF)

For early programming languages such as ALGOL 60, BNF was adequate. However, the increased complexity of modern languages required a more systematic notation to describe their syntax effectively.

Therefore, the so-called Extended Backus-Naur Form (EBNF) was introduced. The inventor of Extended BNF - EBNF, is Niklaus Wirth, a Swiss computer scientist. EBNF is based on BNF but it adds more elements.

“One cannot fail to notice the lack of ‘common denominators.’”

Niklaus Wirth stated this in his first article in the journal “Communications of the ACM” in 1977 about EBNF. [19]

He explained that there is a lack of common standards or unified fundamental principles across different language definitions and notational forms. Each language definition tends to use slightly different variants of the Backus-Naur form based notation. There is no generally accepted standard, and the differences are often small but unnecessarily varied. He stated that a more uniform notation would improve scientific work and comparison. EBNF is his solution for this problem.

As mentioned EBNF is clearly based on BNF, but it adds more elements. EBNF can define exact numbers of elements, describe exceptions, add comments, use flexible rule names, and allow extensions. This makes grammar definitions shorter and it is easier to read them. But EBNF also has restrictions. It can only handle languages with elements in order and is not able to describe very complex grammar. (cf. [16, pp. VII-VIII])

The main disadvantage of BNF according to ISO/IEC 14977 was:

“Backus-Naur Form (used in ALGOL 60) has problems if the metasymbols $<$ $>$ $::=$ occur in the language being defined. Some common forms of construction (e.g. comments) cannot be expressed naturally; other constructions (e.g. repetition) are long winded.” [16, p. VII]

The added advantages, but also the limitations of EBNF, will be discussed in more detail in later chapters.

7.4. Wirth's Approach to Syntax Definition

Having established the historical need for a unified standard in the previous section, it is now necessary to examine the concrete syntactic rules proposed by Niklaus Wirth in 1977. In his article published in the journal *Communications of the ACM* (CACM), Wirth addressed the existing problems of syntactic notation directly. His contribution marked an important step in the formal development of language definitions.

Wirth did not intend to introduce a completely new theoretical concept. Instead, his objective was to improve the existing Backus-Naur Form in a systematic way. BNF had already proven to be useful. However, it contained ambiguities and practical limitations.

The notation presented in his article was therefore designed to be more practical and clearer in structure. It focused on improving readability and standardization rather than expanding theoretical expressive power. The goal was not innovation for its own sake. The goal was clarity, consistency, and easier application in real technical environments.

Wirth listed the following properties to extend BNF:

1. The notation clearly separates meta symbols, terminal symbols, and nonterminal symbols.

Example: `syntax = "production"`.

Here, `syntax` is a nonterminal symbol. `"production"` is a terminal symbol. The equals sign and the dot are structural meta symbols.

2. Symbols used for structure can also be used as normal language symbols. This avoids artificial restrictions.

Example: If quotation marks are used as meta symbols, they can still appear inside text if written twice, like `"` in many programming languages.

3. The notation contains a simple way to express repetition. This reduces the need for recursion in simple cases.
4. The notation avoids explicit representation of the empty string, such as by using `<empty>` or ε .
5. The notation is based on the ASCII character set. This makes it easy to use in most technical environments.

(cf. [19])

In his article, he showed the power of EBNF by defining its own syntax as an example (see figure 2):

“This meta language can therefore conveniently be used to define its own syntax, which may serve here as an example of its use. The word identifier is used to denote nonterminal symbol, and literal stands for terminal symbol. For brevity, identifier and character are not defined in further detail.” [19]

```

syntax      = {production}.
production = identifier "=" expression ".".
expression = term {"|" term}.
term       = factor {factor}.
factor     = identifier | literal | "(" expression ")" |
              "[" expression "]" | "{" expression "}".
literal    = " " " " character {character} " " " " .

```

Figure 2.: Definition of EBNF in EBNF in “What Can We Do about the Unnecessary Diversity of Notation for Syntactic Definitions?” from N. Wirth [19]

7.5. Core Syntactic Extensions

The following overview focuses more in detail on the differences in EBNF and of the Extended Backus-Naur Form (EBNF) as described in Concepts of Programming Languages by Robert W. Sebesta. (cf. [20, pp. 149-151])

It is essential to note that these extensions do not enhance the descriptive power of BNF. The extensions are made to increase the readability and writability of formal grammars. BNF is theoretically sufficient for describing context-free languages, but EBNF provides a more concise notation for complex structures.

7.5.1. Core Syntactic Extensions in EBNF in detail

Sebesta identifies three primary extensions common to most versions of EBNF:

- **Optionality** ([...]):

Brackets are used to denote a part of a right-hand side (RHS) that may occur once or not at all.

Example: A C if-else statement can be described as: `<if_stmt> → if (<logic_expr>) <stmt> [else <stmt>]`.

Instead one uses square brackets to make something optional. For example `["a"]` means “a” or empty.

- **Repetition** ({...}):

Braces indicate that the enclosed part can be repeated any number of times, including zero. This replaces the need for recursion when describing simple lists.

Example: A list of identifiers separated by commas: `<ident_list> → <identifier> {, <identifier>}`.

Example: `{a}` means a, aa, aaa, and so on.

- **Multiple-Choice Options** ((...) and |):

When one element must be selected from a group, the options are grouped in parentheses and separated by the OR operator.

Example: Defining a term with different operators: `<term> → <term> (* | / | %) <factor>`.

In this notation, brackets, braces, and parentheses serve as metasymbols (notational tools) rather than terminal symbols of the language being described. (cf. [20, pp. 149-151])

7.5.2. Further Syntactic Extensions in EBNF based on ISO/IEC 14977

However, there are more differences between BNF and EBNF in the ISO/IEC 14977 definition: (cf. [16, p. VII])

- **Defining an explicit number of items:** In some cases it is needed to describe, that e.g., a field must have an exact length. In BNF, this is difficult. Often, the programmer has explained it then informally in English. Now EBNF solves this by allowing an integer followed by a repetition-symbol (*). With this solution, the programmer can specify the exact number of items. This makes rules much shorter.
- **Exceptions:** Sometimes it would be easier to define a rule by specifying what is not allowed. But BNF cannot easily say “any character except a quote.” EBNF has a solution for that. EBNF uses an except-symbol (-) to define rules by exclusion. This makes definitions shorter and easier. Therefore, for example, a programmer can define a group of symbols and then subtract specific exceptional cases, such as “any character except a quote”.
- **Formal comments:** Comments are an important feature, that BNF is not offering. EBNF provides a formal way to add comments using (* and *). And these notes do not change the machine’s logic.
- **Flexible names:** In BNF, names for categories often need brackets. EBNF uses an explicit comma (,) as a so-called concatenate symbol to link parts of a rule. This allows category names to use multiple words without brackets, meta-identifiers no longer have to consist of a single word or be enclosed in angle brackets. It also ensures that the layout of the syntax like spaces or line breaks has no effect on the defined language.
- **Extensions:** Programmers can add their own features to EBNF. These are called “special sequences” and are placed between question marks (?). This makes the start and end of any extension very easy to find. Extensions were used several times in earlier versions of the solution presented in this thesis, but were eliminated in the final version because there was no longer a need for them.

7.5.3. Comparative Notation BNF and EBNF

While classical BNF is theoretically sufficient for describing any context-free language, it is often complicated and repetitive in practice. The main cause therefore is, that it relies, as described in the previous chapter, exclusively on recursion to express lists and optional elements. In the following illustration one can see the advantage of EBNF introducing specialized meta symbols, such as braces, brackets, and parentheses. [20, p. 153]

BNF and EBNF Versions of an Expression Grammar	
BNF:	
<code><expr></code>	<code>→ <expr> + <term></code>
	<code> <expr> - <term></code>
	<code> <term></code>
<code><term></code>	<code>→ <term> * <factor></code>
	<code> <term> / <factor></code>
	<code> <factor></code>
<code><factor></code>	<code>→ <exp> ** <factor></code>
	<code><exp></code>
<code><exp></code>	<code>→ (<expr>)</code>
	<code> id</code>
EBNF:	
<code><expr></code>	<code>→ <term> { (+ -) <term> }</code>
<code><term></code>	<code>→ <factor> { (* /) <factor> }</code>
<code><factor></code>	<code>→ <exp> { ** <exp> }</code>
<code><exp></code>	<code>→ (<expr>)</code>
	<code> id</code>

Figure 3.: Comparison between BNF and EBNF by Sebesta [20, p. 153]

As demonstrated in the example in figure 3, the structural differences lead to a significant reduction in the number of rules required:

- **Repetition via Braces in this example:** In the BNF version of the expression grammar, an `<expr>` must be defined recursively to handle a sequence of additions or subtractions. EBNF replaces this recursion with braces `{...}`, which denote that the enclosed part can be repeated indefinitely or left out entirely.
- **Multiple-Choice Options in this Example:** The example uses parentheses `(...)` combined with the OR operator `|` to group choices. This for example allows the EBNF version to handle both addition and subtraction in a single line. In BNF it requires a more complex approach.

7.5.4. Modern Variations and Standardization

While an official standard exists (ISO/IEC 14977:1996), it is rarely used in practice. There have several variations of EBNF emerged in recent years, which deviate from the standardized syntax, such as using a colon instead of an equal sign or a dot instead of a semicolon. Also the project described in this thesis partially deviates from ISO standards due to using the EBNF dialect of the Lark Python library.

7.6. ISO/IEC 14977 Standards for Assembler Logic

While the previous chapters describe the advantages of EBNF in terms of readability and compactness, the ISO definition has an additional function. The ISO standard ISO/IEC 14977 also defines strict technical rules.

These are very important for error-free processing by an automated tool such as a script assembler, as they ensure that the grammar definition remains mathematically unambiguous.

7.6.1. Operator Precedence

For every EBNF grammar, the strictly defined hierarchy of meta-symbols is essential, as it prevents ambiguities in the rules. The order of the symbols is clearly defined in the ISO standard without any leeway and is specified as follows: (the strongest is the repetition symbol. This takes precedence, followed by the exclude symbol, and so on).

1. Repetition symbol (*): For specifying exact numbers of repetitions.
2. Except symbol (-): For defining sets by exclusion.
3. Concatenate symbol (,): For explicitly linking consecutive elements.
4. Definition separator symbol (|): For separating alternatives.
5. Defining symbol (=): Formally assigns an expression to a nonterminal.
6. Terminator symbol (;): To clearly mark the end of a production rule.

(cf. [16, p. 1])

7.6.2. Layout Neutrality Through Gap Separators

Ideally, a script assembler should be able to process scripts regardless of their visual formatting. The ISO standard is helpful here, as it defines so-called “gap separators” (such as spaces, tabs, or line breaks) (cf. [16, pp. 5]). These characters have no formal meaning outside of terminal strings. This has the significant advantage of allowing for flexible grammar design without affecting the assembler’s logic. Since gap separators and comments have no formal effect on the defined language, the script assembler can ignore these characters during analysis, which, for example, increases its robustness against different programming styles.

In summary, several advantages arise. The grammar definition remains mathematically consistent for automated tools. Furthermore, it is beneficial that the grammar can be made more readable and understandable for human readers through formatting and annotations.

7.6.3. Syntactic Exceptions

As mentioned earlier, the Except symbol (–) allows the assembler to define specific instruction sets. An example of this would be “all characters except a semicolon.” However, the ISO standard also imposes very clear technical limitations on this function, as otherwise paradoxes could arise:

“If a syntactic exception is permitted to be an arbitrary syntactic factor, Extended BNF could define a wider class of languages than the context-free grammars, including attempts which lead to Russell-like paradoxes... Such a license is undesirable and the form of a syntactic exception is therefore restricted to cases that can be proved to be safe.” [16, p. 2]

A good explanation for the Russell Paradox, formulated by Bertrand Russell, is the so-called barber’s example: In a village, the barber shaves precisely all the men who don’t shave themselves—and only these men. The question of whether the barber shaves himself leads directly to a contradiction: If he shaves himself, he shouldn’t, by definition; if he doesn’t shave himself, he should. (cf. [21])

This example illustrates that such self-referential definitions can lead to logical inconsistencies in formal systems.

7.6.4. Alternative Representations for System Compatibility

To guarantee that a grammar functions on different systems, the ISO standard allows the substitution of certain special characters if they are missing. For example, the characters (`:` and `:`) can replace `{` and `}`, or a period (`.`) can replace a semicolon at the end of a rule. This ensures that the grammar remains readable and functional on different platforms without losing its original structure. (cf. [16, pp. 5])

7.6.5. Special Sequences as a Formal Interface for Assembler Extensions

Although special sequences were already described in the previous section as a general extension option, they are listed again here to specifically clarify their “technical” role as placeholders. In the logic of a script assembler, they serve as a standardized mechanism to mark parts of the language that cannot be defined by EBNF itself (e.g., specific character sets). (cf. [16, p. VII])

By enclosing them in question marks (`?...?`), the grammar remains formally valid for the assembler tool. In this work, they thus can provide the flexibility for possible future functional extensions without violating the fundamental syntactic rules of the ISO standard. (cf. [16, p. 3])

7.6.6. Evaluating the Standard for our Script Assembler

After introducing the theoretical foundations of EBNF, it is necessary to evaluate how the ISO standard was applied in the context of the script assembler. The goal of the project was not to implement a complete parser for the Robot Framework. Instead, the focus was on building a preprocessor for specific control tags such as `@OS`. For this reason, the resulting grammar is not a strict implementation of the ISO standard. It represents a pragmatic combination of formal rules and practical requirements.

A central aspect of this evaluation is the concept of layout neutrality. The ISO standard defines spaces, tabs, and line breaks as so-called gap separators. These elements have no formal meaning and can be ignored by the parser. This principle improves flexibility and readability. However, it could not be applied completely in this project. Inline spaces and tabs between control tokens are treated as insignificant and are therefore ignored. This allows test scripts to remain flexible in their formatting.

Line breaks, however, require a different treatment. Robot Framework scripts are structured line by line. This means that line breaks carry semantic meaning. They define the structure of the script. Therefore, line breaks could not be treated as layout-neutral elements. Instead, they were explicitly defined as structural tokens in the grammar. This is a necessary deviation from the ISO specification to ensure correct processing of the input.

Further deviations from the standard were made in the handling of syntactic constraints. The ISO standard introduces syntactic exceptions to define rules by exclusion. However, this mechanism can become complex and difficult to understand. In some cases, it can even lead to problematic definitions. For this reason, syntactic exceptions were not used in the assembler.

Instead, the implementation relies on regular expressions. Regex provides a practical way to capture and process text patterns. It allows the parser to treat sections of Robot Framework code as plain content. The internal structure of these sections is not analyzed. This simplifies the grammar and reduces its complexity. As a result, the grammar becomes easier to read and maintain.

8. Approaches to EBNF [A]

8.1. Requirements

When creating new formal grammar, one quickly is confronted with a common problem in software development. On one side, you want a clean design, but on the other there is this old legacy system at the company that everyone is used to. For the development of the Script Assembler, this means that it must not only be able to process multi-line structures correctly, but also that the legacy system used at KEBA must continue to function with the new formal grammar, and unfortunately, backward compatibility is a must-have.

The legacy system used before the Script Assembler is a simple preprocessing system, that was based on basic string matching. This solution uses markers that start with an @ sign. Those were already existent:

- **@BRACKETSTART, @BRACKETEND:** These strings marked sections of code. The system just checked if a line exactly matched these strings to know if the following Robot Framework code should be appended at the beginning or the end of the final script.
- **@NO_KETOP:** This was a global option that meant a KeTop device was not needed for the test. It could be written anywhere in the file, and the script just searched the whole document for it without looking at the context.

This solution checked if a specific string was present inside a file or a line and did not use a formal parser. This worked fine for isolated, simple use cases, but not with complex, context-sensitive multi-line structures or dynamic block delineations. Without a parser, the hardcoded logic simply becomes an unmaintainable mess over time.

For this reason, designing the new formal grammar became hard to manage, because three major restrictions had to be taken into account, as clarified in the first meetings by KEBA:

- **Deterministic Parsing:** The system needs strict, unambiguous rules to cleanly differentiate syntactic constructs and reliably deal with multi-line scopes.
- **Backward Compatibility:** It was a strict requirement that the old legacy scripts kept working exactly like before. The functionality had to be preserved. A disruption of the established workflow for the testing team was not an acceptable option.
- **Future Extensibility:** The system must be very robust, because it was stated, that new custom annotations and hardware-specific conditions will definitely be added in the future. A complete rewrite of the parser for every new feature would simply not be an option. .

The testing team at KEBA was already deeply familiar with the `@` prefix. Early design discussions brought up the idea to introduce distinct prefixes to simplify the parsing. The suggestion was to introduce `#` for global options, and `&` for block statements. The idea was, that the tokenization process would become significantly easier and more accurate this way. The evaluation feedback then showed the problem that the distinct prefixes break all existing test suites. A complete disruption of the developers workflow would be the result.

This created a conflict. A new parser is strictly necessary, but the legacy `@` prefix and the overall backward compatibility must still remain intact. This exact conflict dominated the first design phase and brought following solution. Though the strictness of the “Backward Compatibility” restriction was not as clear at that point. To establish a good understanding of the available options and technical trade-offs, two contrasting EBNF approaches were modeled. The evaluation happened before the final syntax design was established.

8.2. Minimal EBNF Approach

One of the two strategies explored was the minimal EBNF approach. The idea behind this specific design was to capture the general script structure as broadly as possible. Theoretically, this offers maximum flexibility and makes future extensions very easy.

Listing 1 demonstrates this broad structural definition. The grammar does not strictly define every possible keyword. Instead, it simply distinguishes between standard Robot Framework lines, single-line operators, and multi-line operators.

Listing 1: The Minimal EBNF grammar

```

1 script = {line, nl}, [line].
2 line = robot | operator | multiline_operator.
3 robot = ? robot code line ?.
4 operator = '@', word, { ws, argument }.
5 multiline_operator = operator, [nl], '{' { robot, nl } '}'.
6 nl = "\n".
7 ws = " " | "\t".
8 argument = ? some argument for example comparison operators or strings ?.

```

8.2.1. Problems of the Minimal Approach

While the parser implementation stays lightweight and adaptable with this approach, there is a blind spot and it completely fails at validating the actual code structure:

1. **Unresolved Ambiguities:** The grammar failed to strictly define and categorize specific keywords. It could not differentiate between standalone options (e.g., `@NO_KETOP`) and block-defining modifiers (e.g., `@VERSION`). The parser was basically blind to the meaning of the words and no correct validation was possible.
2. **Scope Definition Conflicts:** A critical issue emerged during the following design phases when indentation was chosen over explicit brackets (e.g., `{` and `}`) as the primary method for defining scopes. The minimal grammar was just defined too loosely for it to be able to differentiate between either an option that has a body despite it not being allowed to (e.g., developer misinterpreting functionality). The parser was not able to detect the syntax error and invalid code might pass through.

8.3. Detailed EBNF Approach

While the Minimal Approach explored maximum flexibility, a completely different and rather strict grammar was developed. This EBNF approach uses clear rules and a strict separation of all constructs. The main goal was a much better parsing clarity.

As shown in Listing 2, this iteration uses distinct prefixes for different types of constructs. It also introduces explicit termination symbols to make the structure clear.

Listing 2: The Detailed EBNF grammar

```

1  S = { ' ' | '\t' } ;
2  robot_line = S , ?robot_script_line? , S ;
3  option_line = S , '#', S , option , S ;
4  statement_line = S , '&', S , statement , '\n' , { line , '\n' } , S , '&/', S
   , '\n' ;
5  line = robot_line | statement_line | S ;
6  script = { option_line , '\n' } , ( ( { line , '\n' } , [line] ) | [option_line] )
   ;
7
8  statement = version ;
9  option = "NO_KETOP" ;
10 section = "BRACKETSTART" | "BRACKETEND" ;
11
12 version = "VERSION" , S , ( '>' | '<' | '=' | '>=' | '<=' ) , S , ?version_number?
   ;

```

Looking closer at the code, one can see, that in this specific design, `#` was introduced for global options, and `&` was utilized for statements, which were explicitly terminated by `&/`. Since this approach was never fully completed `@` was not used yet, but the idea was to use it only for sections. The grammar also checks the expected arguments very strictly. For instance, the `VERSION` statement forces the user to use mathematical comparison operators and a specific version number format (line 12).

8.3.1. Problems of the Detailed Approach

This approach was very precise and stable. Unfortunately, there were two severe drawbacks that made it unusable for the final product:

1. **Loss of Backward Compatibility:** The introduction of `#` and `&` caused problems with the existing `@` prefix used by the existing option. Changing all the old code was not an option. Furthermore it seemed more intuitive to just have one prefix for all types.
2. **Excessive Rigidity:** The rules for how an argument must look are too strict. For example, the forced mathematical operators in the `version` rule make the grammar very rigid. This means every time a new modifier is added with new arguments, the EBNF has to be updated to match the arguments exactly. This made the Script Assembler harder to extend.

8.4. Conclusion

Testing these two different approaches provided important details for the final EBNF. The Minimal Approach showed the possible dangers of structural ambiguity. It is clear that the parser needs to “understand” the difference between a standalone option and a block-defining modifier. On the other hand, the Detailed Approach proved that

excessive rigidity is also not completely perfect. Making the parsing process too easy for the machine just makes extending unnecessary hard.

The conclusion from this evaluation phase was that the final grammar definitely needed a hybrid design to work. On the inside, it requires the precision of the Detailed Approach, which means in detail that all keywords must be strictly defined and categorized internally so the parser can recognize valid scopes and modifiers correct. But at the same time, is still has to keep the same look and the backward compatibility of the old legacy approach by using the @ prefix, everywhere, while also making the argument structure much more relaxed to make it easier to add extensions later on.

The idea of explicit termination symbols like &/ was completely dropped in favor of indentation-based scoping, and the old “sections” were just changed into standard modifiers.

The combination of all these different requirements finally leads to the final functional grammar used in Script Assembler.

9. A Problem with EBNF [F]

10. Our EBNF [F]

Part IV.

Implementation [A, F, X]

11. Preprocessor [A]

11.1. Introduction

As already established earlier on, the Script Assembler introduces a custom syntax added on top of the Robot Framework. It uses annotations such as `@VERSION` to control the test suit setup. To process this grammar the implementation uses the Lark parsing library (cf. [22]). However while the chosen indentation-based syntax is very intuitive for the developers, it introduces technical difficulties for the parser.

Indentation-based scoping is context-sensitive. To find out if a line belongs to a block, the parser needs to know the indentation level of the previous annotation. That requires it that the current indentation depth is saved. While the Lark library provides its own indenter utility for parsing whitespace-scoped languages, it processes indentation globally (cf. [23]). Because the Script Assembler has to differentiate between its own language and the contained Robot Framework code, Lark's indenter would generate tokens for Robot Framework code indents which might lead to blocks that close too early.

That is why an additional preprocessing step is required before the Lark parser can evaluate the script. This step is handled by a Preprocessor. It's goal is to isolate the meta-language indentation from the content-language indentation before the script is parsed.

11.2. Design and Transformation Rules

11.2.1. Content vs. Meta-Language Conflict

The fundamental problem is that both languages use spaces and indentation to define logic and structure. In the Robot Framework, spaces and tabs separate keywords from arguments, and indentation marks the contents of blocks like `FOR` loops and `IF`

statements. At the same time, the Script Assembler uses indentation to determine which lines belong to a specific modifier. This leads to the problem of having to somehow separate those two kind of indents and thus the need for our own Preprocessor instead of the Lark Indenter.

11.2.2. Explicit Block Scoping

To make the syntax context-free, the implicit visual scoping provided by indentation must be changed to explicit block scoping. Clear start and end markers must be defined. This is comparable to the use of curly braces `{}` in C# or Java, or to opening and closing tags like `<div>` and `</div>` in HTML.

For each modifier found by the Preprocessor, the start of the block is simply the modifier itself. The end of the block, shown by a decrease in indentation in the raw script. The Preprocessor marks this by inserting the keyword `@END`.

With these clear boundaries, the EBNF grammar in Lark is simplified. It only needs to search for the `@END` token to recognize the end of a block, allowing it to completely ignore the surrounding whitespace.

11.2.3. Preserving Indents

The following example uses several annotations nested within standard Robot Framework control structures. Listing 3 shows a typical raw Preprocessor input script written by a test developer.

Listing 3: Raw input script with nested modifiers

```
1  Setup Test Environment
2      @VERSION BIOS == 1.6.1
3          Log    Checking BIOS Version
4          @OS WIN11
5              Install Driver    audio_win11.sys
6          Restart Node
```

In Listing 3, a Robot Framework keyword (`Setup Test Environment`) is the root scope. Within these block, the outer Script Assembler layer performs a BIOS version check. This block contains a normal Robot command (`Log`) followed by an inner modifier that checks the operating system. This inner modifier contains a command for a driver installation. Finally, a `Restart Node` command is executed within the BIOS block, but outside the inner operating system check.

The Preprocessor should calculate the active layers of modifiers and remove only the leading indents belonging to these specific layers. Listing 4 shows the expected output.

Listing 4: Expected Preprocessor output

```
1
2 Setup Test Environment
3 @VERSION BIOS == 1.6.1
4     Log     Checking BIOS Version
5 @OS WIN11
6     Install Driver     audio_win11.sys
7 @END
8     Restart Node
9 @END
```

Comparing the two listings shows perfectly how the Preprocessor is supposed to work. The modifier lines themselves (lines 2 and 4 in Listing 4) are completely stripped of their indentation. This is needed since the EBNF grammar only sees modifiers that begin with an @ symbol, so if it is padded with tabs it is not recognized.

Additionally and most importantly the indentations of the contained language are preserved while the modifier indents are not. For example in the raw text the line `Install Driver audio_win11.sys` was indented three times. Once for the BIOS version check modifier, once for the OS check modifier and once for the actual Robot Framework block. In the result the two modifier indents are removed, leaving only the Robot Framework one. `Restart Node`

Finally, the decrease of indentation before the `Restart Node` command triggers the Preprocessor to insert the @END tag and the open end adds another one lines 6 and 8 in Listing 4).

11.3. Evolution of the Algorithm

Developing an appropriate algorithm for this was a technical challenge. The algorithm has to correctly identify mixed indentation levels and clearly distinguish between a line that defines a new lexical scope and a line that just belongs to an existing scope. Parsing context-sensitive, whitespace-dependent meta-languages is a highly complex process, as demonstrated by the evolution of the Preprocessor logic.

11.3.1. Linear Depth Counting Approach

The first version of the Preprocessor used a simple linear mechanism for counting layers. The basic idea was that the depth of a modifier block directly correlated to a simple integer counter.

In this naive approach, the algorithm maintained a single integer variable called `current_layer`, starting at zero. The Preprocessor iterated through the script line-by-line, searching for the `@` symbol. When it found a modifier, it incremented the `current_layer` counter by one. For all following lines, the algorithm removed a number of leading indents exactly equal to the value of `current_layer`.

To find the end of a block, the algorithm compared the total number of indents on the current line against the `current_layer` variable. If a line had fewer indentations than the active layer count the algorithm assumed that the block had ended. It then inserted an `@END` tag and decremented the counter. This logic worked in simple situations. However, it contained a fundamental flaw regarding relative and absolute indentation depths.

11.3.2. Collision with the Structure of the Content

The linear depth counting approach unfortunately failed when Script Assembler annotations were mixed with Robot Framework commands that use indents inside a modifier. At that point the absolute count of tabs of a line is no longer just determined by the Script Assembler modifier depth.

Listing 5 demonstrates the specific edge case that exposed the flaw in the naive algorithm.

Listing 5: Problematic edge case

```

1  @VERSION BIOS == 1.2.3
2      FOR      ${item}      IN      @{ITEMS}
3          @VERSION KBDriver != 0.2.5
4              Add Item To Cart      ${item}
5          Sleep      0.5s
6      END

```

When verifying the algorithm with Listing 5, the error occurs in line 5.

1. **First Modifier:** On line 1, the first `@VERSION` modifier sets the `current_layer` to 1.

2. **Inner Modifier:** On line 3, the inner `@VERSION` modifier increments the `current_layer` to 2.
3. **Conflicting Line:** On line 5, the `Sleep` command has exactly 2 leading tabs (one from the outer modifier, and one from the `FOR` loop).

The algorithm compares the `Sleep` line's tab count (2) to the `current_layer` (2). Since the numbers are equal, the algorithm assumes that the `Sleep` command is still in the inner `@VERSION` block. But this is wrong, since the inner `@VERSION` block started with two tabs, thus requiring any code inside of it to have 3 tabs. This means that the `Sleep` command actually exits the inner block.

That is why the algorithm unfortunately fails, as seen in Listing 6. It is counting tabs in `current_layer` integer for the indentation and completely ignores how many of these tabs were actually introduced by the contained language.

Listing 6: Wrong output of the problematic script with the first approach

```

1  @VERSION BIOS == 1.2.3
2  FOR    ${item}    IN    @{ITEMS}
3  @VERSION KBDriver != 0.2.5
4      Add Item To Cart    ${item}
5  Sleep    0.5s
6  @END
7  END
8  @END

```

11.3.3. Stack-Based Approach

To fix this collision, the algorithm required a different approach. Instead of tracking a simple depth integer, it had to explicitly store the absolute indentation level at the exact moment each modifier scope was opened.

The revised Preprocessor uses a stack to keep track of the indents. Whenever the Preprocessor finds a new modifier, it calculates the exact number of leading tabs on that line and pushes this absolute tab count onto the top of the stack.

This approach maps the nested structure of the text to a linear state history. The value at the top of the stack is always the baseline tab count of the currently active innermost modifier.

Re-evaluating Listing 5 using this stack-based approach yields the correct behavior:

- **Line 1:** Outer modifier has 0 tabs. Stack becomes: [0].
- **Line 3:** Inner modifier has 2 tabs. Stack becomes: [0, 2].

- **Line 4:** Content has 3 tabs. The value 3 is greater than 2 (top of stack). The content belongs to the inner block.
- **Line 5:** The `Sleep` command has 2 tabs. The value 2 is less than or equal to 2 (top of stack).

Listing 7: Correct output of the problematic script with the final approach

```

1  @VERSION BIOS == 1.2.3
2  FOR    ${item}    IN    @{ITEMS}
3  @VERSION KBDriver != 0.2.5
4      Add Item To Cart    ${item}
5  @END
6      Sleep    0.5s
7  END
8  @END

```

When the algorithm processes line 5, it sees that the current tab count (2) is less than or equal to the top of the stack (2). This triggers the scope resolution logic. The Preprocessor removes the top value from the stack and correctly inserts the `@END` tag to terminate the inner `@VERSION` block. Finally, it reaches the end of the file, and since it still has one modifier to close, it ends the output with a final `@END`.

By using a stack, the Preprocessor can perfectly keep track of at which level a modifier starts and also where it should end. This makes it entirely unaffected of whatever indentation structures the content uses.

11.4. Implementation Details

The theoretical concepts and the stack-based state machine were implemented in Python. The resulting `Preprocessor` class encapsulates the state variables, the regular expressions for token identification, and the transpilation loop that processes the script line-by-line.

11.4.1. Initialization and Utilities

The Preprocessor is initialized with specific regular expressions. These expressions identify valid modifiers and are passed to the Preprocessor by the Parser, removing the need to maintain them in two places. Furthermore, the class defines utility functions to manage the indentations. Listing 8 shows this class initialization alongside these core utility functions.

Listing 8: Class initialization and indentation utility functions

```

1  import re
2  from broker.managers.case.script_assembler.parser_exception import
    ParserSyntaxError
3  from logging import Logger
4
5  log: Logger = Logger("SCRIPT ASSEMBLER")
6
7  class Preprocessor:
8      modifier_regex: str
9      option_regex: str
10     indent_regex: str = r"^(?:\t| {4})+"
11
12     def __init__(self, modifier_regex: str, option_regex: str) -> None:
13         self.modifier_regex = modifier_regex
14         self.option_regex = option_regex
15
16     def count_indents(self, line: str) -> int:
17         """Counts the number of indents at the start of a line."""
18
19         match = re.match(self.indent_regex, line)
20         if not match:
21             return 0
22
23         indent_str = match.group()
24
25         # Count number of indents
26         tabs: int = indent_str.count('\t')
27         spaces: int = indent_str.count(' ')
28
29         return tabs + spaces
30
31     def remove_indents(self, line: str, amount: int) -> str:
32         """Removes 'amount' indents from the beginning of a line."""
33
34         match = re.match(self.indent_regex, line)
35         if not match:
36             return line # No indents to remove
37
38         indent_str = match.group()
39         indents = re.findall(r"\t| {4}", indent_str)
40
41         # Remove x indents
42         remaining_indents = ''.join(indents[amount:])
43         rest_of_line = line[len(indent_str):]
44
45         return remaining_indents + rest_of_line

```

An important detail in Listing 8 is the definition of `indent_regex`. Different code editors and IDEs format files differently. The regular expression `r"^(?:\t| {4})+"` handles this reliably. It matches any continuous sequence of leading whitespace. It strictly treats either a single tab character or exactly four spaces as one logical unit of indentation.

This regular expression is then used by two utility functions. First, the `count_indents` method calculates the indentation depth of a line. It sums the `\t` characters and blocks of four spaces. This process reduces the visual indentation into a single integer.

Second, the `remove_indents` method removes the tabs of the meta-language while preserving the tabs of the content. It stores all indentations into an array. After that, it removes the first elements based on the `amount` parameter. This parameter corresponds to the current depth of the modifier stack. Finally, it restores the line by puts back together the remaining indents with the actual code.

11.4.2. Main Processing Loop

The `preprocess` method handles the line-by-line transformation. Listing 9 shows the initialization of the state tracking variables: the `result` string buffer, the `layer_tabs` stack, and the `is_modifier_empty` flag.

Listing 9: Initialization of the main preprocessing loop

```

1
2     def preprocess(self, script: str) -> str:
3
4         log.info("Starting preprocessing of robot script...")
5         result: str = ""
6         layer_tabs: list[int] = []
7         is_modifier_empty: bool = False
8
9         for line in script.splitlines():
10             stripped: str = line.strip()
11             if stripped == "":
12                 continue
13
14             log.debug(f"Processing robot script line: \"{stripped}\"")
15
16             mod: bool = False
17
18             if stripped.startswith("@"):
19                 if stripped == "@END":
20                     raise ParserSyntaxError("Cannot have @END in script since it
21                                             is an internal keyword")
22                 mod = re.fullmatch(self.modifier_regex, stripped.split('
23                                     ')[0][1:]) is not None
24
25             # [...]
```

First, the loop skips any empty lines. Before using the slower `re.fullmatch` operation, the algorithm checks if the line starts with an `@` symbol and thus is a modifier. This filters out the majority of lines. If the line begins with that symbol, we check if the modifier is `@END`. Since this is an internal token only used for terminating the blocks, users are not allowed to use it in raw scripts. That means if it exists in the code, we raise a `ParserSyntaxError`.

11.4.3. Scope Handling

After reading a line, the Preprocessor decides whether the line belongs to the current scope, opens a new scope, or closes existing scopes. Listing 10 illustrates this stack management logic.

Listing 10: The stack implementation for resolving indentation scopes

```

1
2     # [...]
3     if len(layer_tabs) > 0:
4         tabs: int = layer_tabs[-1]
5         cnt: int = self.count_indents(line)
6
7         if mod:
8             line = stripped # dont want any indents on modifier lines
9
10        if cnt <= tabs:
```

```

11         while len(layer_tabs) > 0 and cnt <= layer_tabs[-1]:
12             layer_tabs.pop()
13             result += "@END\n"
14
15             if is_modifier_empty:
16                 raise ParserSyntaxError("Cannot have empty modifiers")
17
18             result += self.remove_indents(line, len(layer_tabs)) + "\n"
19
20             if mod:
21                 layer_tabs.append(cnt) # adds count of tabs to the stack
22                 log.debug(f"Processing modifier; Increasing layer by one")
23                 is_modifier_empty = False
24
25             else:
26                 result += f"{line}\n"
27
28             # [...]

```

If the stack is not empty, the algorithm compares the absolute indentation count of the current line (`cnt`) with the top value of the stack (`tabs`). If the current count is less than or equal to the top value, a scope ends.

This resolution is implemented using a `while` loop rather than a simple `if` statement. This ensures proper handling when a single line with fewer indents closes several nested modifier blocks simultaneously. For each level removed, the string `@END\n` is appended to the result.

The empty modifier validation occurs inside this loop. If a block is closed but the `is_modifier_empty` flag is still set to `True` because no inner content was processed, transpilation aborts. This is needed since a modifier without code to modify is pointless.

That is why the Preprocessor assumes that it is most likely either not correctly indented or a typo by the test developer and needs to raise an `ParserSyntaxError` to communicate that. Finally, the `remove_indents` utility is called to strip the indents of the meta-language before the line is appended to the final output stream.

11.4.4. Finalization

At the end of the loop, new scopes are initialized, and the output is verified before being passed to the Parser.

Listing 11: Handling the root scope and finalizing the output stream at the end of the file.

```

1
2         # [...]
3
4         if mod:
5             is_modifier_empty = True
6
7             if len(layer_tabs) == 0:

```



```
8
9             # If there aren't any previous modifiers, add initial root
              layer
10             layer_tabs.append(0)
11
12
13         # Close any remaining open scopes
14         if len(layer_tabs) != 0:
15
16             for _ in range(0, len(layer_tabs)):
17                 result += "@END\n"
18
19         log.info("Finished preprocessing the robot script!")
20
21         return result.replace('\r', '')
```

If the current line is a modifier and the stack is empty, a root level of 0 is pushed to the stack to serve as the absolute baseline.

After iterating through the entire file, the algorithm does a final cleanup. If the file ends without returning to zero indentations, the stack still contains unclosed modifier blocks. The logic iterates through all remaining items in the `layer_tabs` stack, adding an `@END` tag for each. Finally, carriage returns are normalized to avoid cross-platform compatibility issues.

11.5. Summary

The Preprocessor acts as a hidden translation layer between the human-readable script and the more machine-centric requirements of the Parser. By using a stack the algorithm properly stores and processes the indentations of the meta-language. It removes the assembler's scoping indicators while retaining the precise indentation required for the Robot Framework runtime environment. Furthermore, the Preprocessor also handles the first validation by rejecting manually added `@END` keywords and empty modifiers.

12. Parser [F, A]

13. Syntax Highlighting [A]

13.1. Introduction

One part of this project presented in this thesis, was the development of Syntax Highlighting. As already described in previous chapters, the Script Assembler works by overlaying a custom syntax on top of the standard Robot Framework language using an EBNF grammar, thus allowing annotations like `@VERSION` or `@BRACKETSTART`.

There is a fundamental problem though. Because standard text-editors and native Robot Framework syntax highlighters do not recognize this new grammar, custom annotations would be displayed as plain text, or worse, marked as syntax errors.

For KEBA, the frontend editor needs a way to visually mark these custom elements, because if they are not able to do so, it can lead to worse developer experience and more errors that go unnoticed until execution. Without visual validation, developers make more syntax mistakes. For example, writing `@VERISON` instead of `@VERSION`) or using an invalid argument type, like passing a plain string where a semantic version number is expected, would not be recognized while typing. These mistakes remain unnoticed. The errors are only found during the backend transpilation phase much later, when an error log is generated. This breaks the entire programming workflow and wastes precious time.

To fix this, custom syntax highlighting should be integrated into the web-based pre-existing code editor. The main objective simply was to get immediate visual feedback and validation. Correct keywords, different types of arguments and errors should all be highlighted in unique colors. This should reduce the time needed for test development using the annotations by a lot. Figure 4 shows how the editor is supposed to handle valid and invalid inputs.

```
1
2  @VERSION BIOS <= 1.0.0
3  |   ... Execute on KeTop ... echo VALID
4  |
5  @VERSION BIOS <= invalid
6  |   ... Execute on KeTop ... echo INVALID
7  |
8  @VERSION BIOS EQUALS 1.0.0
9  |   ... Execute on KeTop ... echo INVALID
10 |
11 @VERISON BIOS <= 1.0.0
12 |   ... ^TYPO => no syntax highlighting
13
```

Figure 4.: Syntax Highlighting in Action

13.2. Monaco Editor Setup

The frontend application already used the Monaco Editor to provide basic syntax highlighting for standard Robot Framework scripts with the Monarch library. (cf. [24, 25]). Monaco is the underlying technology of Visual Studio Code and can be used in frontends through packages. The package that was used was outdated and a few other architectural issues had to be fixed, before the syntax highlighting could be implemented.

13.2.1. Legacy Version Conflicts

Over time the frontend ecosystem of the project has been updated to Angular 15, but the existing editor integration was still based on the widely used wrapper library `ngx-monaco-editor`. Unfortunately, this specific package was no longer being updated. It officially only supported Angular up to version 12 (cf. [26]). Previously developers used the npm force flag during installation to make the project compile using this already deprecated package, bypassing the dependency conflicts.

This workaround made it possible for the legacy code to somehow continue working, but it also created major risks for maintainability and the build process. Forcing npm installation is a bad practice since it only covers up package mismatches and other errors. Because of this, there is a major risk that one might get unpredictable runtime

behaviors (cf. [27]). Since we do not want that, especially later on in production, the package needed to be replaced.

After some research, the best and most widely used package seemed to be `ngx-monaco-editor-v2`. This version is an actively maintained community fork that is compatible with modern Angular versions. (cf. [28]).

The migration was handled in several steps to ensure a stable base. The `package.json` had to be updated, then the `app.module.ts` was refactored to match the new library and finally the asset paths in `angular.json` had to be updated.

13.2.2. Custom Theme

A custom Monaco theme was required to map the custom syntax to the correct colors. The requirement for that theme, was for the colors to stand out compared to the standard Robot Framework syntax.

The editor worked fine with its predefined default designs, but as soon as a newly created custom design was applied, a runtime error occurred: `Uncaught TypeError: this.themeData.colors is undefined`. This error completely stopped the editor from rendering any custom syntax highlighting.

Looking in the documentation gave me a hint in the right direction. The engine of Monaco is rather strict, so a `colors` object has to be provided for a theme, even if it is empty and never used. (cf. [29]). The solution was to add an empty `colors` object to the theme definition. Listing 12 shows the working theme configuration, which includes the required `colors` property alongside the specific token rules.

Listing 12: Custom monaco theme including the required color object

```
1 // Definition of the custom language theme
2 (<any>window).monaco.editor.defineTheme('robot-theme', {
3   base: 'vs-dark',
4   inherit: true,
5   rules: [
6     { token: 'variables', foreground: '#9CDCFE' },
7     { token: 'keywords', foreground: '#ffe18f', fontStyle: 'bold' },
8     { token: 'modifier', foreground: '#f09c00', fontStyle: 'bold' },
9     { token: 'invalid', foreground: '#ff0000', fontStyle: 'underline italic' },
10    // [...]
11  ],
12  colors: {} // Required to prevent 'this.themeData.colors is undefined'
13 });
```

13.3. Monarch Lexer for Custom Annotations

After the editor and theme were successfully configured, the core task was to define the syntax rules for the custom grammar of the Script Assembler. The already in use Monaco Editor uses Monarch for syntax highlighting, which works declaratively. Developers define lexers using JSON configuration objects instead of writing complex parser code (cf. [25]).

13.3.1. Tokenization and State Machine

Monarch basically just uses Regex to handle the tokenization. The step-by-step-process starts with the lexer, that reads the document text in order and tries to match the defined regular expression patterns. When it finds a match, the matched text is assigned a specific token class. This can be a specific keyword or a comment or it could also just be a newly defined custom modifier. The editor then maps these token classes to the specific color values defined in the active theme.

A global list of Regex rules is sufficient for simple languages. But for the custom annotations there is a need for context-aware parsing. For example, the arguments that follow an `@VERSION` annotation should be highlighted, but if the same argument is written in a comment it should not.

Monarch's state machine solved this problem. The lexer starts in a default `@root` state, when a specific pattern is matched, the lexer jumps into a different, explicitly defined state, and different Regex rules apply in this new state.

13.3.2. Handling Modifiers and Arguments

The annotations can be roughly divided into two categories. Those without any arguments and those with specific argument structures.

Modifiers without arguments, such as `@BRACKETSTART`, are simple to implement. As shown in Listing 13, when the lexer encounters this keyword, it transitions to an `@endOfLine` state. This state makes sure that no more text is accepted on this line.

Listing 13: Monarch tokenizers root state

```
1  tokenizer: {  
2    root: [  
3      // Standard Robot Framework variables  
4      [/[\\$@]\\{[\\w_\\/]+\\}/, 'variables'],  
5    ]  
  }
```

```

6      // #region MODIFIERS
7      [/@(?::PREAMBLE)\b/, 'modifier', '@preambleArg'],
8      // Modifiers without any arguments transition to @endOfLine
9      [/@(?::BRACKETSTART|BRACKETEND)\b/, 'modifier', '@endOfLine'],
10     // Modifiers with complex arguments transition to their specific states
11     [/@(?::VERSION)\b/, 'modifier', '@versionArgsIdentifier'],
12     [/@(?::OS)\b/, 'modifier', '@osArg'],
13     // #endregion
14
15     // [...]
16 ],
17 // ...

```

Modifiers with arguments are more difficult, because they require a more complex sequence of state transitions. The modifier `@VERSION` is a good example of this as the version check needs three different components. Also important is, they have to be in an exact order, starting with an identifier (the BIOS or software name), then a comparison operator (e.g., `<`, `>=`), and at last a valid semantic version string.

A chain of states was implemented, Root to identifier, Identifier to operator, Operator to version. That means in detail, when `@VERSION` is found in the `@root` state, the lexer transitions to `versionArgsIdentifier`. Once a valid word is found there, it transitions to `versionArgsOperator` and finally to `versionArgsSemver`.

This “state machine” approach guarantees that the syntax highlighting exactly fits to the strict sequence that the parser needs. Listing 14 shows this state chain in detail.

Listing 14: State transitions for the `@VERSION` modifier

```

1  versionArgsIdentifier: [
2      // Entry point: match a word, then proceed to operator state
3      [/[w_]+/, 'arg.word', '@versionArgsOperator'],
4      [/\n/, 'invalid', '@root'],
5      [./, 'invalid']
6  ],
7  versionArgsOperator: [
8      // Match comparison operators, then proceed to semantic version state
9      [/(?:[<>]=?)|(?:[!=]=)/, 'arg.operator', '@versionArgsSemver'],
10     [/\n/, 'invalid', '@root'],
11     [./, 'invalid']
12 ],
13 versionArgsSemver: [
14     // Match standard semantic versioning format
15     //
16     https://semver.org/#is-there-a-suggested-regular-expression-regex-to-check-a-semver-string
17     [/(0|[1-9]\d*)\.(0|[1-9]\d*)\.(0|[1-9]\d*)(?:-((?:0|[1-9]\d*|\d*[a-zA-Z-][0-9a-zA-Z-]*)?(?:\.[0-9a-zA-Z-]*)?)/, 'arg.version'],
18     [/\n/, 'invalid', '@root'],
19     [./, 'invalid']
20 ],

```

The Regex for the semantic version argument might look complicated and hard to maintain, but it is the officially provided Regex by the organization behind semantic versions.

13.4. Troubleshooting and Error Handling

After solving the correct highlighting of valid syntax, the next tasks were to handle edge cases and wrong inputs.

During the implementation of the lexer, two main problems appeared, firstly to prevent “state bleed” over multiple lines, secondly to provide visual error feedback without using a dedicated Language Server.

13.4.1. Fixing State Bleed

The custom EBNF grammar for the Script Assembler is strictly line-bound, so an annotation has to be strictly on a singular line. So if a line break occurs the editor should return to the standard context used for the syntax highlighting of the Robot Framework.

In the beginning, this caused a big conflict with Monarch’s stateful tokenization. When the lexer goes into a state like `versionArgsIdentifier`, it just stays inside it until an explicit command such as `@root` or `@pop` is triggered. If a developer typed the keyword `@VERSION` and then pressed the Enter key without any arguments, the lexer got stuck in the `versionArgsIdentifier` state on the next line. Because of this, normal Robot Framework code on the following lines was wrongly checked against the version argument regular expressions. This completely destroyed the syntax highlighting for the whole rest of the file.

To fix this “state bleed”, the lexer had to recognize line breaks as termination triggers. Normally, Monarch works line-by-line and does not pass line break characters to the Regex engine (cf. [30, pp. 8-9]). To change this, the `includeLF` property was enabled in the language configuration (see Listing 15).

Listing 15: `includeLF` in the Monarch config

```
1  (<any>window).monaco.languages.setMonarchTokensProvider('robot', {
2    ignoreCase: true,
3    includeLF: true, // Exposes newline to the regex engine
4    tokenizer: {
5      // tokenizer rules ...
6    }
7  });
```

With `includeLF` set to true, line breaks can be recognized. As can be seen in the state machine definition in Listing 14, each argument state received a new rule for handling the newline character: `[/\n/, 'invalid', '@root']`. Sometimes the lexer just resets

to the `@root` state without marking an error. This is used when a line break is allowed and does not count as a syntax mistake in that specific case. This ensures that a line break always returns the lexer to the `@root` state and so protects the highlighting of all following code.

13.4.2. Visualizing Syntax Errors

A major goal for the fronted-editor was to show syntax errors to the user right away, but there was a strict rule for this. The error highlighting had to work only through the Monarch configuration, because complex language server or an extra linter were not an option, as tools like that just use too many resources.

This was quite a problem because Monarch is really just a lexical analyzer, meaning a lexer, and not a semantic parser. (cf. [25]; [31, pp. 5-8]). It only uses regular expressions to map strings to token classes. It does not know anything about “errors” or “warnings”. It also cannot draw the typical red wavy lines, that are used for code defects in modern IDEs.

To solve this problem, a visual workaround had to be built with a custom token class called `invalid`. In the custom Monaco Theme (see Listing 12), this category got a unique style with red text-color, italic style and an underline.

This `invalid` category then served as a catch-all mechanism in the state machine. Monarch always checks the rules of a state from top to bottom (cf. [25]). By placing the rule `[./, 'invalid']` at the very end of a states rule list, any character that does not match the expected valid patterns or the newline reset automatically falls through to this fallback rule.

A good example for this is the modifier `@BRACKETSTART`. As it does not allow any arguments, the following happens by running the system. When the lexer sees this exact keyword, it switches to the `endOfLine` state (see Listing 16).

Listing 16: The `endOfLine` state

```
1 endOfLine: [  
2   [/\n/, '', '@root'], // Valid: line ends, return to root  
3   [./, 'invalid']      // Error: any other character is invalid  
4 ],
```

For example for the case, when a developer accidentally types a `@BRACKETSTART` argument, the lexer first recognizes the annotation, jumps to the `endOfLine` state. There, it expects the end of the line. Instead, it encounters a space, and the word


“argument”. Because these characters do not match the `/\n/` Regex, they fall through to the `/./` rule. As a result, the string “argument” turns red, gets underlined and italic, and is signaling by this way a syntax error to the developer. The user sees a clear syntax error right away without backend validation. Real-time error feedback is a useful feature, and even better, everything remains strictly within the boundaries of the Monarch lexer architecture.

13.5. Limitations

Being able to catch errors early with the syntax highlighting was definitely worth the addition. But even with all effort, there is a limitation that could not be fixed, since there is a conflict between the rules of the two used languages, which could only be fixed through an own Language Server.

The EBNF grammar for the backend is strictly case-sensitive. That means, all annotations and modifiers like `@VERSION` or `@OS`, have to be in uppercase, and different spelling leads to a syntax error during transpilation.

The standard Robot Framework language on the other hand works differently. Here upper and lower case do not matter., the language is case-insensitive (cf. [32]). Because of this, the `ignoreCase` property of the Monarch lexer must be set to true globally, if not, the normal keywords and variables do not get the correct highlighting (see Figure 5 for example).



```
1
2 Setup Test Environment
3 ... Execute on KeTop ... echo Syntax Highlighting still works
4
5 @VERSION BIOS <= 1.1.1
6 ... Execute on Ketop ... State Bleed
7
```

Figure 5.: State bleed in the editor when line breaks are not included

The Monarch API only allows this setting for the entire language, because it is only possible through the global `ignoreCase` property defined in the `IMonarchLanguage` interface. A custom override for one tokenizer rule(`IMonarchLanguageRule`) or a single state is not possible. Therefore, the setting affects everything, the entire tokenizer. In the implemented editor configuration this global setting also affect the custom annotation

rules, and as a result, if a developer types `@version` in lowercase, the frontend lexer still matches the rule and applies the corresponding syntax highlighting. But with the same script, the validation crashes when the backend EBNF parser takes the script. The reason is that the parser strictly needs the uppercase syntax(cf. [33]).

This limitation leads to a small difference between the visual representation and the backend, but it is an unavoidable compromise. Disabling `ignoreCase` globally completely breaks the standard Robot Framework code highlighting, and hat specific code is the vast majority of the file. The more important thing is, that the base language simply has to function. This was the basis for the final architectural decision together with KEBA.

14. Extensibility [F]

15. Database [X]

Besides the extension of the Script Assembler's capabilities and the creation of a new EBNF language that makes the software more extendable, another main focus of the project is the creation of a mockup of the TestRail database. This mockup is not only made for testing purposes, but allows the user to work independently of the TestRail instance. The mockup is designed to be as close to the real TestRail database as possible, allowing a seamless switch between the two instances. This means that all functionalities that are available in the real TestRail database are also available in the new mockup.

This section aims to give an overview of the different types of databases, the structure and functionalities of the TestRail database specifically, the resulting design and implementation of the mockup database, and the way the database switch is finally implemented.

15.1. Types of Databases

Depending on the requirements of a given project, different types of databases are suitable for different use cases. It is important to analyze what qualities are needed and which qualities are negligible to find the best fitting solution.

The most important qualities of a database type can be expressed in the following categories:

- **Data Model:**

The data model of a certain database type describes how the data is structured and stored. This can be in the form of tables, documents, graphs or other structures. The data model has a notable impact on aspects like performance and scalability.

- **Storage Architecture:**

The storage architecture of a database type not only describes how the data is

stored, but also how it may be accessed, managed and distributed. This can be in the form of a single server, a cluster of servers or even a distributed network of servers across the globe.

- **Query Model:**

The query model of a database type describes the means of accessing the data. This can occur in the form of a specifically designed query language, an API or other methods. This has an undeniable impact on the ease of use of the database and how the user writes and optimizes queries.

Based on these qualities, various databases types have been developed, each with their own unique advantages and disadvantages. Some of the most common types of databases include but are not limited to relational databases, object-oriented databases and NoSQL databases, which can be further categorized by the data model they use, such as document stores, key-value stores, column-family stores and graph databases. (cf. [34])

15.2. Relational Databases

For the purpose of this project, we will take a closer look at relational databases, as this is the type of database that TestRail uses.

15.2.1. Structure of Relational Databases

Relational databases are based on the relational model, which organizes its data into tables, consisting of rows and columns. Each table represents a specific entity and each row in the table represents an instance of that entity. The columns in the table represent the attributes of the entity.

Relationships between these entities can be established through the use of keys. A primary key is a unique identifier for each instance of an entity, while a foreign key is a reference to the primary key of another entity. Meaningful relationships between entities can be built with the usage of these keys, allowing complex queries to be executed across multiple linked entities.

Grouping in queries allows for the aggregation of data based on specific attributes, while joins allow for the combination of data from multiple tables based on the established relationships.




Figure 6.: Simple Diagram showcasing the relationship between three entities in a relational database.

In figure 6, we can see a simplified example of a relational database structure. The three entities "Artist", "Label" and "Album" are represented as tables, with their respective attributes as columns. The relationships between the entities are established through the usage of primary and foreign keys, allowing for queries to be executed across multiple linked entities. For example, we can query for all albums by a specific artist or a specific label.

15.2.2. ACID

ACID is an acronym that stands for the four key properties of a safe database transaction:

- **Atomicity:**

Atomicity ensures that a transaction, which is often a sequence of multiple operations, is treated as a single unit. This means that either all operations within a given transaction are executed successfully, or none of them are. If any operation of the transaction fails due to an error, a power outage or any other issue, the entire transaction is rolled back to its previous state before the transaction began. This guarantees that the database remains consistent and prevents data corruption or loss that could occur if only part of the transaction were to be executed.

- **Consistency:**

Consistency ensures that the database can only be brought from one valid consistent state to another, while following all defined rules, constraints and triggers. Invalid transactions that violate any of the defined rules are rolled back accordingly, ensuring that the database remains consistent.

- **Isolation:**

Isolation ensures that concurrent transactions have no impact on each other. This means that the effects of a transaction in progress are not visible to other transactions until it completes successfully. This prevents a number of issues that would otherwise arise from concurrent transactions, like dirty reads, where one transaction reads data that has been modified by another transaction but not yet fully committed, potentially leading to incorrect results.

- **Durability:**

Durability ensures that once a transaction has been successfully completed and committed, the state of the database persists even in the case of a crash, power failure or any other issue.

Databases that adhere to these properties are considered to be reliable and safe for use in critical environments, as they guarantee that the data not only remains consistent, but also that it is not lost or corrupted due to unforeseen events. (cf. [35])

15.2.3. Advantages and Disadvantages of Relational Databases

One of the main strengths of relational databases is their ability to handle long and complex queries across multiple linked entities with the usage of foreign keys and joins. The consistency and integrity of the data is guaranteed by the ACID properties. On top of that, relational databases are some of the most widely used and well-known types of databases, which means that there are a lot of resources and a large community available for support and development.

However, relational databases can struggle with scalability and performance when dealing with large volumes of data or high traffic. The fixed schema of relational databases can also be a disadvantage, as it can make it difficult to adapt to changing requirements or to handle unstructured data. Additionally, the execution of complex queries can lead to performance issues on larger datasets.

15.3. The TestRail Database

TestRail is a web-based test management tool that provides a wide range of features for managing and organizing test cases, test runs and test results. The relational database in the background of TestRail is used to store all of the data related to these features. The user can interact with the database directly by using the TestRail API, although the database itself is not directly accessible to the user.

Thus, for the creation of the mockup database, the structure and functionalities of the TestRail database have to be studied and analyzed based on the API documentation and the structure of its response data. While the mockup database doesn't have to be an exact replica of the TestRail database, the functionalities used by the Script Assembler need to deliver responses with identical structure to ensure the switch between databases can occur without any issues.

In order to achieve this, the most important entities and their relationships have to be identified.

15.3.1. Test Cases

Test cases can be seen as the most fundamental entity in the TestRail system, as they represent the individual tests that are to be executed. Test cases have a number of attributes, such as a title, description, type and the test steps that need to be executed. Test cases are furthermore linked to test suites and sections.

15.3.2. Projects, Suites and Sections

Projects are the highest level of organization in TestRail, representing a specific testing effort or initiative. Each project can contain multiple test suites, which are collections of test cases. The test cases of Suites can further be organized into sections, allowing for more specific grouping and categorization. These relationships are an essential part of the database structure that needs to be replicated.

15.3.3. Test Runs

A test run represents a specific execution of test cases. Test runs are linked to test cases and contain information about the execution, such as the status of each case, the tester who executed the run and any relevant comments or other metadata. Test runs make

the tracking and management of test executions possible. Importantly, a test run can only be created for test cases that are contained in one singular test suite.

15.3.4. Test Plans

Test plans are used to organize and manage test runs. TestRail allows the user to group multiple test runs into a test plan, which alleviates the need to manage each run individually.

15.3.5. Test Results

Test results represent the outcome of a test case within a test run. Test results mainly consist of a status, which can contain, but is not limited to:

- Passed
- Failed
- Untested
- Retest
- Blocked

15.3.6. Users

Users represent the individuals who interact with the TestRail system. Users can have different roles and permissions within the system. They can be assigned to test runs, test plans or other entities as creators. With the user entity it is possible to track who is responsible for which run and plan.

15.4. PostgreSQL

For the implementation of the mockup database, PostgreSQL was chosen as the database management system, as it is an open-source database that is not only widely used but also well-known for its reliability and performance.

PostgreSQL is not only a database management system that the team is familiar with, but it also provides all the necessary features and capabilities needed for the

implementation of this part of the project. It supports the relational data model, which is needed for replacing the TestRail database, and it guarantees the reliability and consistency of its data with the ACID properties.

15.4.1. pgAdmin

The open-source management tool "pgAdmin" is used for the administration of the PostgreSQL database, as it provides an easily navigable user interface for management, query execution and other administrative tasks.




Figure 7.: Screenshot of the pgAdmin interface. (cf. [36])

The current version of pgAdmin, which is pgAdmin 4, provides a wide range of features for managing and administering PostgreSQL databases, some of which include but are not limited to the following: (cf. [36])

- **SQL query tool:**

Inside pgAdmin, there is a built-in SQL editor that allows users to write and execute SQL queries directly against the database. This tool provides syntax highlighting, further aiding the user in writing correct queries.

- **Direct data editing:**

The tables in a database can be directly edited through the interface, allowing for quick and easy modifications of single points of data without the need to write and execute SQL queries. Caution should be taken when using this feature however, as the quick and easy nature of this feature can lead to unintended consequences.

- **Management of multiple databases:**

The pgAdmin interface allows the management of multiple databases within the same interface, allowing for intuitive switching between independent databases.

15.5. First Approaches

Early into the development of the database, two main approaches were considered for how the database should interact with the backend of the main system.

One approach was to create an API that would allow the main system to interact with the database through API calls. This would have resulted in a clear separation between the database and the main system.

The other approach was to work in the preexisting backend to connect to the database directly. This approach was ultimately chosen, as it not only makes the switch between the two databases simpler, but also integrates the functionality of the new database directly into the main system. This approach was further endorsed by the supervisor of the project, as it helps the system to eventually transition away from the TestRail database and towards the new mockup database, which allows for more independence and expandability in the future.

The python library "psycopg" is used for the connection between the main system and the PostgreSQL database. This library provides a simple and efficient way to connect to the PostgreSQL database and execute various commands and queries. It also provides support for connection pooling, which can help to improve the performance of the database interactions.

15.6. Database Connections in the System

The following section aims to give a brief look into the structure and functionalities of not only the existing TestRail database connection, but also the new PostgreSQL mockup database. The goal is to show how the new database is designed to replicate the implementations of the TestRail connection, while also highlighting clear differences on a technical level.

15.6.1. Existing TestRail Functionality

The existing functionalities concerning the TestRail connection is contained within the TRSession class, which is responsible for all interactions with the database. This class contains functions for retrieving and creating test cases, runs, plans and results, as well as other relevant data.

The communication with the TestRail database is achieved through the usage of the TestRail API, which provides all necessary reading and writing functionalities.

```
def get_case_title(self, case_id: int, safe_name: bool = False) -> str | None:
    """
    Returns the Name of a test case by a given test case id
    Parameters:
        case_id (int):    Name (Title) of the testcase. (partial match with contains)
        safe_name (bool): if true: delete the ":" character in the title for file system compliance
    Returns:
        (string): The ID of the test case in section: section_name containing: testcase_name
        None: no test case was found
    """

    response = self.do_request(f"get_case/{case_id}", TRConst.REQUEST_TYPE_GET)
    if not response:
        return None
    if len(response) <= 0:
        self.log.warning(f"could not find a result for test case: {case_id}")
        return None
    if not safe_name:
        return response['title']
    return response['title'].replace(":", "")
```

Figure 8.: Code snippet of the get_case_title function within the TRSession class.

In figure 8, one of the functions of the TRSession class can be seen. This function is responsible for retrieving the title of a test case based on its ID. Using the ID of the wanted test case, a GET request is sent to the TestRail API, which then returns the relevant data. The function then processes the response data and returns the title.

All functions of the TRSession class follow a similar structure, with the type of interaction with the API being the main difference. For example, functions that create new entities in the database use POST requests and send the relevant data in the body of the request.

15.6.2. DBSession

The DBSession class is designed to replicate all used functionalities of the existing TRSession class, while replacing the communication with the TestRail API with a direct connection to the PostgreSQL database, using the "psycopg" python library. It is strictly necessary that the functions of the DBSession have the same name, parameters

and return values as the functions of the TRSession class, in order to avoid major conflicts when switching between the two database connections.

```
def get_case_title(self, case_id: int, safe_name: bool = False) -> str | None:
    """
    Returns the Name of a test case by a given test case id
    Parameters:
        case_id (int):    Name (Title) of the testcase. (partial match with contains)
        safe_name (bool): if true: delete the ":" character in the title for file system compliance
    Returns:
        (str | None): on success: title of found case - else: None
    """

    if case_id == None:
        print("Did not get a case_id")
        return None

    if case_id < 0:
        print(f"Invalid value for case_id: {case_id}")
        return None

    self.cursor.execute("SELECT title FROM test_cases WHERE id = %s", (case_id,))

    result = self.cursor.fetchone()

    if result != None:
        return str(result[0])

    return None
```

Figure 9.: Code snippet of the `get_case_title` function within the `DBSession` class.

In figure 9, the same function as in figure 8 can be seen, but this time within the `DBSession` class. Thus, while the parameters and return value of the function are identical to the one in the `TRSession` class, the implementation is different.

Instead of interacting with an API, this function uses a connection to the PostgreSQL database to execute a SQL query that retrieves the title of the specified test case.

15.7. Integration of the DBSession Class

Since the `DBSession` class replicates all essential functions of the `TRSession` class, the switch between database usages can be achieved easily.

Any occurrence of the `TRSession` class within the backend of the system is now denoted as a `TRSession` or `DBSession`, using a pipe symbol. This means that at system startup, the appropriate session class has to be identified and initialized accordingly.

This way, the `TRSession` functionalities remain unaltered, while the possibility of using the `DBSession` is a new feature that can be easily switched on or off.

```
def fetch_cases(self, testrail: TRSession | DBSession):  
    """ updates the cases by reading the cases from Testrail  
    """
```

Figure 10.: Code snippet showcasing a function that takes a parameter that can be either a TRSession or a DBSession.

15.7.1. Switching Between Databases

The switch between the two databases is implemented in a way that allows the user to easily and definitively choose which database should be used at the execution of the python script. This is achieved by using a command line argument that can be passed when executing the script. The argument is then processed and the appropriate database instance is initialized based on if the argument is present or not.

```
python ./broker.py
```

Here, the standard TestRail database is used, as no argument is passed to the script.

```
python ./broker.py -db  
# OR  
python ./broker.py --database
```

Here, the PostgreSQL mockup database is used, as the argument is passed to the script. The argument can either be a short version, or a long version, as both are processed in the same way and lead to the same result.

15.8. The PostgreSQL Database

Since the structure of the TestRail database is not publicly available, the structure of the PostgreSQL mockup database is based on the structure of the response data the TestRail API provides. All entities contain their respective attributes, as well as keys that establish the relationships between the different entities.

It is important to note that the relationships between entities are not made with the usage of foreign keys, as the logic for the relationships is implemented in the backend of the system, and the database is only used for storing data. While this is not the most efficient way to implement the relationships between entities, as a mockup database, it is sufficient for the purpose of this project. Despite that, in section 15.9, potential

future steps are discussed, one of which is the implementation of proper relationships between entities with the usage of foreign keys.

15.9. Potential Future Steps

15.9.1. Easier Creation of Test Cases

TestRail provides a functionality for the creation of test cases through the API, as well as through their user interface. This makes the creation of new test cases in the database an easy task, especially with the user interface, as the user can efficiently create new test cases and other entities without having to worry about the structure of the database itself.

Currently, the creation of test cases in the PostgreSQL database is done manually by either executing SQL queries manually or by making use of the pgAdmin user interface. This is not only time-consuming, but also prone to errors, as the user has to ensure that all necessary attributes of the test case are filled in correctly and that the relationships between entities are established properly.

A potential future step to fix this issue would be to create a user interface that allows for the easier creation of test cases and other entities in the database. This system could either be a separate application that interacts with the database on its own, or it could be integrated into the existing system.

15.9.2. Encryption

For convenience and ease of use, the users of the PostgreSQL database currently have their credentials stored in plain text within their respective columns in the database. While this is not an issue for the mockup database, as it is only used for testing purposes, if the PostgreSQL database approach is to be used in the future as a replacement for the TestRail database, it would be necessary to implement a system for encrypting the credentials of all users, in order to ensure the security of the database and the data it contains.

This could be achieved by using a hashing algorithm to hash the passwords before storing them in the database, and then comparing the hashed password with the stored

hash when a user attempts to log in. This would ensure that actual passwords cannot be retrieved from the database, even if the database is compromised.

15.9.3. Implementation of Proper Relationships

As mentioned in the previous section, the relationships between entities in the PostgreSQL database are not implemented with the usage of foreign keys, but rather with the usage of keys that are processed in the backend of the system. This is not the most efficient way to implement relationships, as it can lead to issues with data integrity, as well as performance issues when executing queries that involve multiple linked tables.

The logical next step to improve the implementation of the database would be to implement proper relationships between entities with the correct usage of foreign keys.

Glossary

EBNF Extended Backus–Naur form

IDE Integrated Development Environment

TRP TestRail Procedure

1

References

- [1] Python Software Foundation, “What is Python? Executive Summary,” retrieved 2026-02-21. [Online]. Available: <https://www.python.org/doc/essays/blurb/>
- [2] Parewa Labs Pvt. Ltd., “Python Modules (With Examples),” retrieved 2026-03-01. [Online]. Available: <https://www.programiz.com/python-programming/modules>
- [3] Python Software Foundation, “Installing Python Modules,” retrieved 2026-03-01. [Online]. Available: <https://docs.python.org/3/installing/index.html>
- [4] Erez Shinan, “Lark documentation,” retrieved 2026-03-01. [Online]. Available: <https://lark-parser.readthedocs.io/en/latest/>
- [5] Raphaël Barrois, “semantic-version · PyPI,” retrieved 2026-03-01. [Online]. Available: <https://pypi.org/project/semantic-version/>
- [6] Tom Preston-Werner, “Semantic Versioning 2.0.0,” retrieved 2026-03-01. [Online]. Available: <https://semver.org/>
- [7] Daniele Varrazzo, “The Psycopg 3 project,” retrieved 2026-03-01. [Online]. Available: <https://www.psycopg.org/psycopg3/>
- [8] Robot Framework Foundation, “Robot Framework,” retrieved 2026-03-01. [Online]. Available: <https://robotframework.org/>
- [9] Microsoft, “Monaco Editor,” retrieved 2026-03-01. [Online]. Available: <https://microsoft.github.io/monaco-editor/>
- [10] The PostgreSQL Global Development Group, “PostgreSQL: About,” retrieved 2026-03-01. [Online]. Available: <https://www.postgresql.org/about/>
- [11] Nikhil Garg, “What is a Keyword in Programming?” retrieved 2026-03-01. [Online]. Available: <https://www.geeksforgeeks.org/dsa/what-is-keyword-in-programming/>
- [12] Robot Framework Guides, “Style Guide,” retrieved 2026-03-01. [Online]. Available: https://docs.robotframework.org/docs/style_guide

- [13] Python Software Foundation, “Indentation in Python,” retrieved 2026-03-01. [Online]. Available: https://docs.python.org/3/reference/lexical_analysis.html#indentation
- [14] Elly Johnson, “Pros and Cons of Robot Framework,” retrieved 2026-03-01. [Online]. Available: <https://testautomationtools.dev/pros-and-cons-of-robot-framework-general-overview/>
- [15] Sonal Tuteja, “Introduction to Recursion,” retrieved 2026-03-02. [Online]. Available: <https://www.geeksforgeeks.org/dsa/introduction-to-recursion-2/>
- [16] ISO/IEC 14977, “Information technology - Syntactic metalanguage - Extended BNF,” *International Standard 14977:1996(E)*, vol. 1996, pp. 1–24, 1996. [Online]. Available: http://www.iso.org/iso/catalogue_detail.htm?csnumber=26153
- [17] Wikipedia contributors, “Metalanguage,” 2026, accessed: 2026-02-23. [Online]. Available: <https://en.wikipedia.org/wiki/Metalanguage>
- [18] H. Mössenböck, *Compiler Construction*, 2025.
- [19] N. Wirth, “What Can We Do about the Unnecessary Diversity of Notation for Syntactic Definitions?” vol. 20, no. November, pp. 1–2, 1977.
- [20] R. W. Sebesta, *Concepts of Programming Languages ELEVENTH edition*.
- [21] Wikipedia contributors, “Barbier Paradoxon,” 2026, accessed: 2026-02-23. [Online]. Available: <https://de.wikipedia.org/wiki/Barbier-Paradoxon>
- [22] Lark, “Lark on github,” 2026, accessed: 2026-03-01. [Online]. Available: <https://github.com/lark-parser/lark>
- [23] —, “Lark docs: Indented tree,” 2026, accessed: 2026-03-01. [Online]. Available: https://lark-parser.readthedocs.io/en/latest/examples/indented_tree.html
- [24] Microsoft, “Monaco editor,” 2026, accessed: 2026-02-27. [Online]. Available: <https://github.com/microsoft/monaco-editor>
- [25] —, “Monarch syntax highlighting documentation,” 2026, accessed: 2026-02-27. [Online]. Available: <https://microsoft.github.io/monaco-editor/monarch.html>
- [26] atularen, “ngx-monaco-editor: Monaco editor component for angular 2 and above.” 2024, accessed: 2026-02-24. [Online]. Available: <https://www.npmjs.com/package/ngx-monaco-editor>

- [27] npm, Inc., “npm config: force,” 2026, accessed: 2026-02-27. [Online]. Available: <https://docs.npmjs.com/cli/v11/using-npm/config#force>
- [28] miki995, “ngx-monaco-editor-v2: Monaco editor component for angular 2 and above,” 2026, npm package registry entry, accessed 2026-02-27. [Online]. Available: <https://www.npmjs.com/package/ngx-monaco-editor-v2>
- [29] Microsoft, “Monaco editor api: IStandaloneThemeData.colors,” 2026, accessed: 2026-02-27. [Online]. Available: https://microsoft.github.io/monaco-editor/typedoc/interfaces/editor_editor_api.editor.IStandaloneThemeData.html#colors
- [30] T. van Krieken, “Deriving Syntax Highlighting Grammars,” Master’s thesis, Eindhoven University of Technology, 2023. [Online]. Available: <https://homepages.cwi.nl/~jurgenv/theses/TarVanKrieken.pdf>
- [31] A. Aho, M. Lam *et al.*, *Compilers: Principles, Techniques, and Tools*, 2nd ed. Boston: Addison-Wesley, 2006.
- [32] Robot Framework Foundation, “Robot framework user guide,” 2026, section 3.3.2 User Keyword Names; Accessed 2026-02-27. [Online]. Available: https://robotframework.org/robotframework-RFCP-syllabus/docs/chapter-03/user_keyword#332-user-keyword-names
- [33] Microsoft, “monaco.d.ts: Monaco editor typescript api definitions (v0.34.1),” 2022, accessed: 2026-02-27. [Online]. Available: <https://unpkg.com/monaco-editor@0.34.1/monaco.d.ts>
- [34] Vinayak Baranwal, “What are the different types of databases? explained with use cases and architectures,” 2025, accessed: 2026-02-28. [Online]. Available: <https://www.digitalocean.com/community/conceptual-articles/database-types>
- [35] Wikipedia contributors, “Acid,” 2026, accessed: 2026-02-28. [Online]. Available: <https://en.wikipedia.org/wiki/ACID>
- [36] pgAdmin Development Team, “pgadmin documentation - user interface,” 2026, accessed: 2026-03-02. [Online]. Available: https://www.pgadmin.org/docs/pgadmin4/latest/user_interface.html

List of Figures

1.	Simple Diagram depicting the order scripts with a certain modifier keyword are run in.	19
2.	Definition of EBNF in EBNF in “What Can We Do about the Unnecessary Diversity of Notation for Syntactic Definitions?” from N. Wirth [19] . .	25
3.	Comparison between BNF and EBNF by Sebesta [20, p. 153]	28
4.	Syntax Highlighting in Action	53
5.	State bleed in the editor when line breaks are not included	59
6.	Simple Diagram showcasing the relationship between three entities in a relational database.	64
7.	Screenshot of the pgAdmin interface. (cf. [36])	68
8.	Code snippet of the get_case_title function within the TRSession class.	70
9.	Code snippet of the get_case_title function within the DBSession class.	71
10.	Code snippet showcasing a function that takes a parameter that can be either a TRSession or a DBSession.	72

List of Tables

List of Listings

1.	The Minimal EBNF grammar	35
2.	The Detailed EBNF grammar	35
3.	Raw input script with nested modifiers	42
4.	Expected Preprocessor output	43
5.	Problematic edge case	44
6.	Wrong output of the problematic script with the first approach	45
7.	Correct output of the problematic script with the final approach	46
8.	Class initialization and indentation utility functions	46
9.	Initialization of the main preprocessing loop	48
10.	The stack implementation for resolving indentation scopes	48
11.	Handling the root scope and finalizing the output stream at the end of the file.	49
12.	Custom monaco theme including the required color object	54
13.	Monarch tokenizers root state	55
14.	State transitions for the @VERSION modifier	56
15.	includeLF in the Monarch config	57
16.	The endOfLine state	58

Appendix